

Introduction To Haskell Programming

S. P. Suresh

Chennai Mathematical Institute

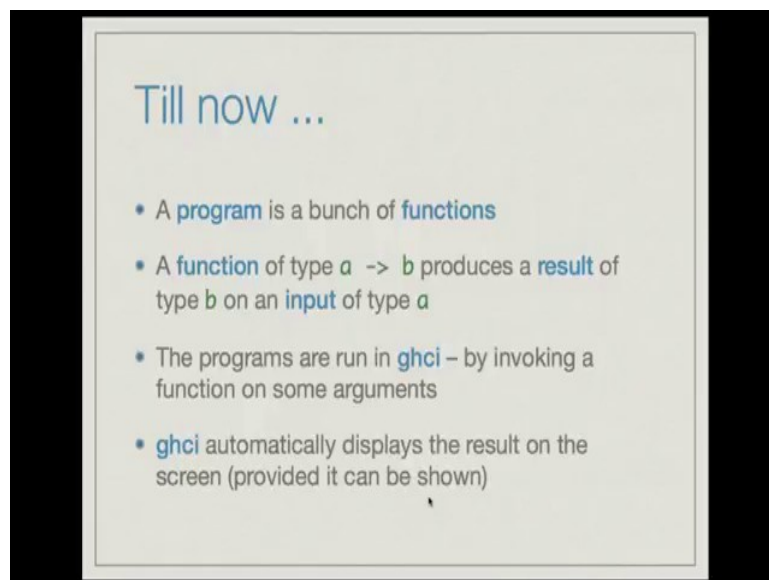
Module # 07

Lecture - 02

Input / Output

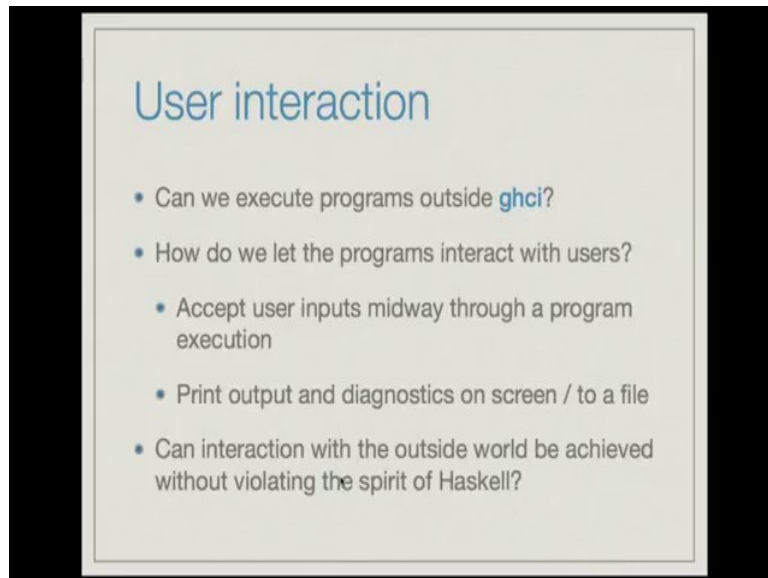
In this lecture, we shall study Input Output in Haskell.

(Refer Slide Time: 00:08)



Till now, the view we have taken is that a program is a bunch of functions, a function of type $a \rightarrow b$ produces a result of type b on an input of type a . The programs are run in `ghci` by invoking a function on some arguments, `ghci` automatically displays the result on the screen provided the result can be shown or in other words, provided the result is a type that belongs to the type class `Show`.

(Refer Slide Time: 00:46)




User interaction

- Can we execute programs outside `ghci`?
- How do we let the programs interact with users?
 - Accept user inputs midway through a program execution
 - Print output and diagnostics on screen / to a file
- Can interaction with the outside world be achieved without violating the spirit of Haskell?

But, this is a limited form of interaction with the user, we would like a slightly better user interaction model like in other programming languages. So, the questions we have are , can we execute programs outside `ghci`, how do you let the programs interact with users that is accept user inputs midway through a program execution. Print output and diagnostics on a screen or to a file, can interaction with the outside world be achieved without violating the spirit of Haskell. So, these are some of the questions that we considered in this lecture.

(Refer Slide Time: 01:25)

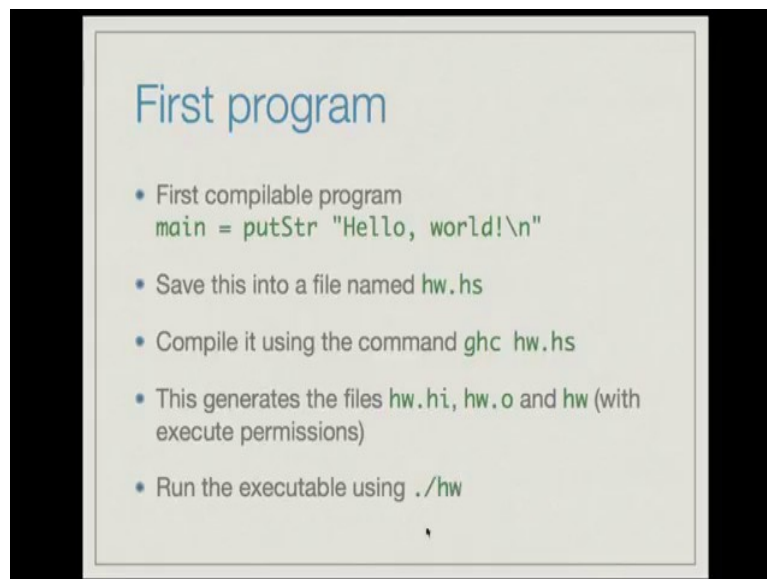


Standalone programs and `main`

- Execution of a Haskell program starts with the function `main`
- Every standalone Haskell program should have a `main` function

We shall first consider standalone programs, execution of a Haskell program starts with the function `main`, this is not something we have seen till now. Till now, our program was just a bunch of functions and `ghci` automatically interpreted any function that we entered in it. But, if you have to write a standalone program, execution has to start at some place and that is the function `main`, every standalone Haskell program should have a `main` function.

(Refer Slide Time: 02:03)



The slide is titled "First program" in a blue font. It contains a bulleted list of instructions for creating and running a Haskell program. The first bullet point includes a code snippet for the `main` function. The subsequent bullet points describe saving the file, compiling it with `ghc`, the resulting files, and running the executable with `./hw`.

- First compilable program

```
main = putStr "Hello, world!\n"
```
- Save this into a file named `hw.hs`
- Compile it using the command `ghc hw.hs`
- This generates the files `hw.hi`, `hw.o` and `hw` (with execute permissions)
- Run the executable using `./hw`

Here is an example program, in fact this is the simplest compilable program that we can write, `main = putStr "Hello, world! \n"`. So, this as the name suggests puts a string on the screen, so the way we run this is to save this in to a file named `hw.hs`, `hw` standing for Hello world. We compile this file using the command `ghc hw.hs`. Compiling this generates the following files, `hw.hi`, `hw.o` and `hw`, the file of interest was the `hw` without any extension.

If you see this on a Unix terminal, we will see that the permissions for these two are read write, but `hw` has execute permissions or in other words `7 5 5`. We can run the executable using `./hw`, dot denotes the current directory, dot slash `hw` means that the executable files `hw` can be found in the current directory. If you add the path, where your Haskell executable resides in your path environment, you can just invoke the functions using `hw`, but for now in the examples we will just use dot slash programming.

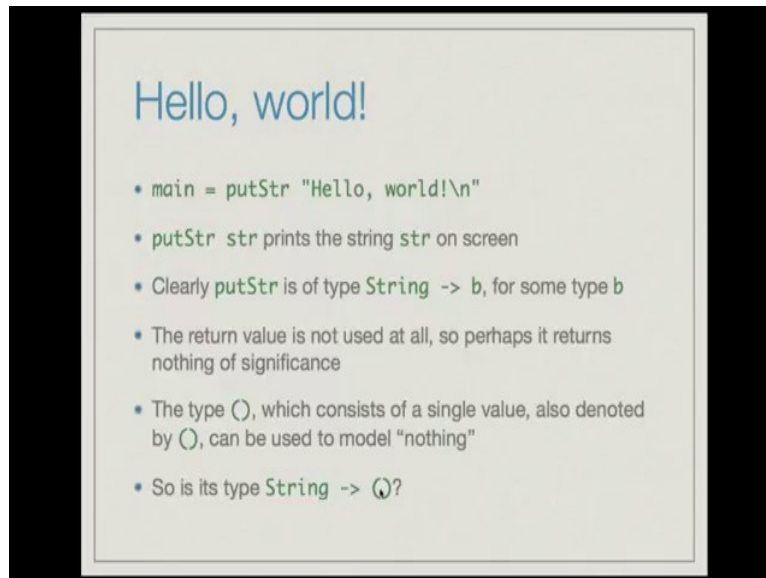
(Refer Slide Time: 03:45)



GHC is the Glasgow Haskell Compiler, `ghci` that we have been using till now is the interactive version of the compiler, one can view `ghci` as an interpreter or a play ground in which you test your programs. If the program is intended for use by others, then it is usually written as a standalone program, compiled using `ghc` and shipped. Compiled versions of programs run much faster and use much less memory as compared to running them in `ghci`.

You can check the various options that `ghc` offers by typing `ghc --help` in the terminal, you can use `ghc --show-options` to know all the options that you can provide, all the compiler options that you may provide to `ghc`, but this is a huge list. If you want to know more about `ghc` and `ghci`, you can consult the GHC Manual at this url which is the part of the official Haskell page. So, we have learnt how to write a simple program which stands alone and runs on its own, compile it and run it.

(Refer Slide Time: 05:05)

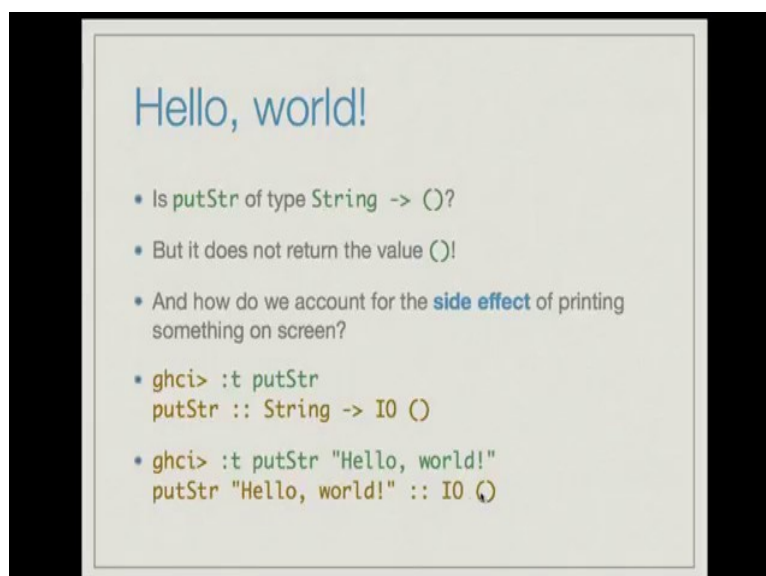


Hello, world!

- `main = putStr "Hello, world!\n"`
- `putStr str` prints the string `str` on screen
- Clearly `putStr` is of type `String -> b`, for some type `b`
- The return value is not used at all, so perhaps it returns nothing of significance
- The type `()`, which consists of a single value, also denoted by `()`, can be used to model "nothing"
- So is its type `String -> ()`?

Let us study the program in more detail, `main = putStr "Hello, world! \n"`, `putStr` is a function, the behavior is that `putStr str` prints the string `str` on screen. So, clearly `putStr` is of type `String -> b` for some `b`, because the input is a string, but you notice that in this main program the return value is not used at all. So, perhaps we can say that `putStr` does not return anything of significance, the type `()` or empty as it is called denotes nothing or it can be used to model nothing. This type empty is denoted like an empty tuple and it consists of a single value, which is also denoted by an empty tuple, so the question is, is the type of `putStr`, `String -> ()` ?.

(Refer Slide Time: 06:25)

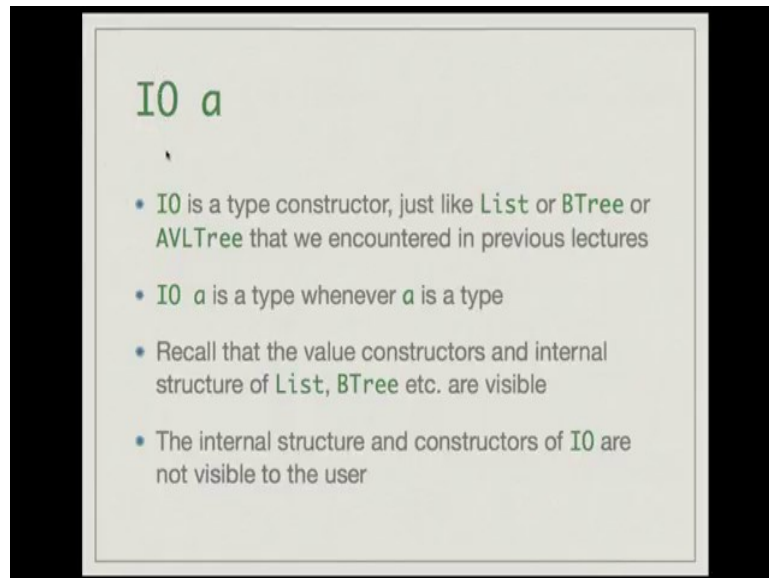


Hello, world!

- Is `putStr` of type `String -> ()`?
- But it does not return the value `()`!
- And how do we account for the **side effect** of printing something on screen?
- `ghci> :t putStr`
`putStr :: String -> IO ()`
- `ghci> :t putStr "Hello, world!"`
`putStr "Hello, world!" :: IO ()`

But, we notice that `putStr` is not an expression that returns a value and more over, it has a side effect, which is that of printing something on screen. So how do we account for the side effect? So, if you actually type `:t putStr` in `ghci`, you will see that `ghci` says what the type of `putStr` is, the type is `String -> IO ()`. If you check what the type of `putStr "Hello world!"` is, you will see that, the type is `IO ()`.

(Refer Slide Time: 07:13)

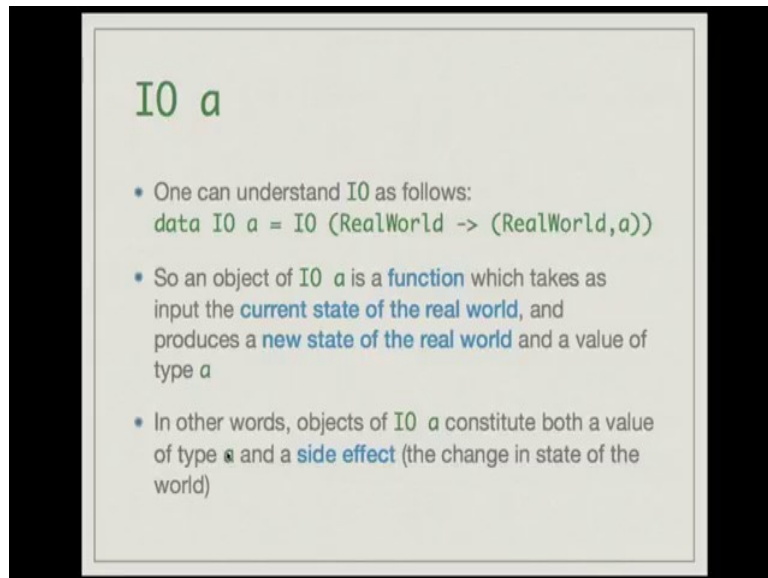


IO a

- `IO` is a type constructor, just like `List` or `BTree` or `AVLTree` that we encountered in previous lectures
- `IO a` is a type whenever `a` is a type
- Recall that the value constructors and internal structure of `List`, `BTree` etc. are visible
- The internal structure and constructors of `IO` are not visible to the user

So, what is this `IO`? `IO` is a type constructor, just like some other type constructors we have encountered in previous lectures like `List` or `BTree` or `AVLTree`, etc. So, therefore, `IO a` is a type whenever `a` is a type, but there is a distinction. Recall that the value constructors and internal structure of user defined data types like `List`, `BTree` etc are visible. But, the internal structure and constructors of `IO` are not visible to the user. In other words the user cannot do any kind of pattern matching based on the constructors of `IO`.

(Refer Slide Time: 07:57)



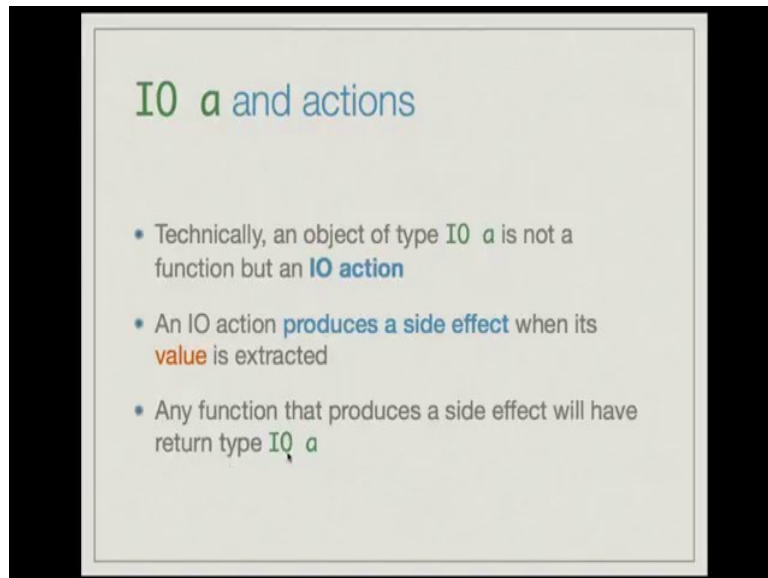
IO a

- One can understand IO as follows:
`data IO a = IO (RealWorld -> (RealWorld,a))`
- So an object of **IO a** is a **function** which takes as input the **current state of the real world**, and produces a **new state of the real world** and a value of type **a**
- In other words, objects of **IO a** constitute both a value of type **a** and a **side effect** (the change in state of the world)

One way to understand IO is as follows ,its declaration can be thought of as `data IO a = IO (RealWorld -> (RealWorld,a))` , the IO on the left is the type constructor, the IO on the right is the value constructor. And the values are of type `RealWorld -> (RealWorld,a)`. RealWorld here is not an actual Haskell type, but it is just one way we can understand what IO means. Assume that there is a type, which represents all states of the real world.

So, we can think of IO as taking as input the current state of the real world and producing a new state of real world due to side effects and also producing the value of type **a**. In other words object of IO a constitute both a value of type **a** and a side effect namely the change in state of the world.

(Refer Slide Time: 09:06)

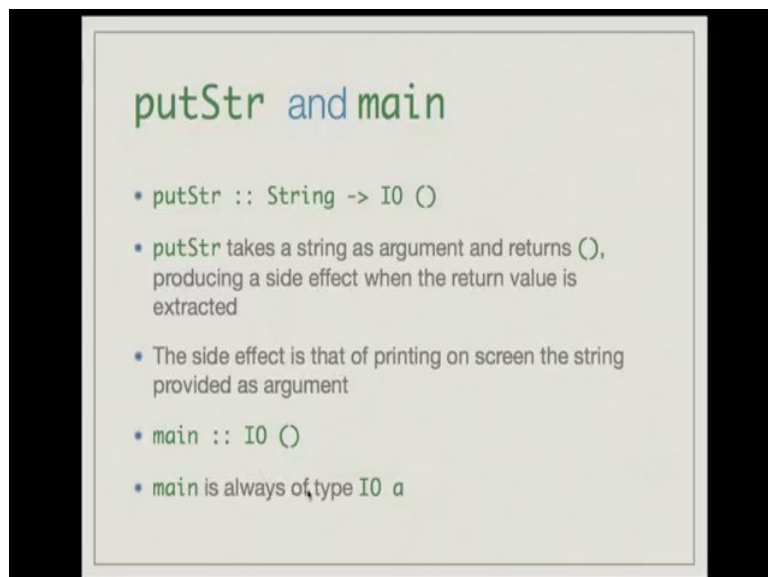


IO a and actions

- Technically, an object of type `IO a` is not a function but an **IO action**
- An IO action **produces a side effect** when its **value** is extracted
- Any function that produces a side effect will have return type `IO a`

Technically an object of type `IO a` is not usually referred to as a function, but as an IO action. This is an important distinction. An IO action produces a side effect when its value is extracted. Any function that produces a side effect will have return type `IO a`.

(Refer Slide Time: 09:32)



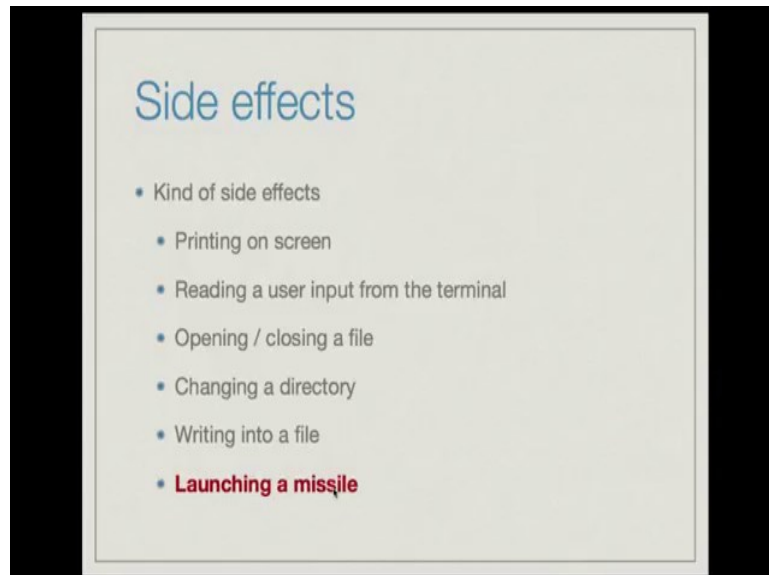
putStr and main

- `putStr :: String -> IO ()`
- `putStr` takes a string as argument and returns `()`, producing a side effect when the return value is extracted
- The side effect is that of printing on screen the string provided as argument
- `main :: IO ()`
- `main` is always of type `IO a`

So, let's get back to `putStr`. `putStr` has a type `String -> IO ()`. `putStr` takes a `String` as the argument and returns an empty tuple `()`. And in the process of producing the empty tuple as output it also produces a side effect when the return value is extracted. The side effect in the

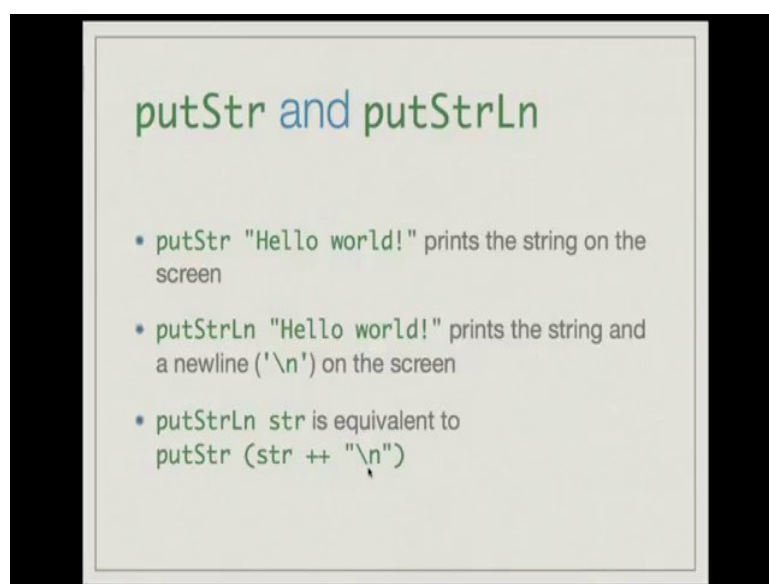
case of `putStr` is that of printing on screen the string that is provided as argument. `main` as you can see is of type `IO empty`. `main` is always of type `IO a` for some `a`.

(Refer Slide Time: 10:15)



Now, we talked about side effects in the course of executing an action, what kind of side effects can happen. Examples are printing on screen, reading the user input from the terminal, opening or closing a file, changing the directory, writing into a file etc or may be launching a missile, driving a truck etcetera.

(Refer Slide Time: 10:45)



Now, there is a close variant of `putStr` namely `putStrLn`, you can read it as put String line. `putStr "Hello world!"` prints the string on screen, whereas `putStrLn "Hello world!"` prints the string and appends a new line on the screen. So, `putStrLn str` is equivalent to `putStr (str ++ "\n")`

(Refer Slide Time: 11:17)



The slide is titled "Chaining actions" in blue text. It contains four bullet points and a code snippet. The first bullet point states that the `do` command is used to chain multiple actions. The second bullet point shows a code snippet: `main = do` followed by two indented lines, `putStrLn "Hello!"` and `putStrLn "What's your name?"`. The third bullet point explains that `do` makes actions take effect in sequential order, one after the other. The fourth bullet point states that indentation is important.

Chaining actions

- We use the command `do` to chain multiple actions
- ```
main = do
 putStrLn "Hello!"
 putStrLn "What's your name?"
```
- `do` makes the actions take effect in **sequential order**, one after the other
- Indentation is important

Stand alone actions are not of much use unless you can perform a lot of actions and Haskell provides a way to chain actions. We use the command `do` to chain multiple actions. For example, you could say

```
main = do
 putStrLn "Hello!"
 putStrLn "what's your name?"
```

`do` makes the actions take effect in sequential order one after the other in the order presented in the program text. Here the indentation is important in the `do` command, but since the indentation is sometimes hard to keep track of...

(Refer Slide Time: 12:00)



Haskell offers an alternative friendlier syntax, which is to use braces and semi colons.

```
main = do {
 putStrLn "Hello!";
 putStrLn "what's your name?";
}
```

And these actions can also occur inside let, where etc. For instance,

```
main = do {act1; act2;}
 where
 act1 = putStrLn "Hello, ";
 act2 = putStrLn "world!"
```

So, you can define local actions.

(Refer Slide Time: 12:42)

## More actions

- `print :: Show a => a -> IO ()`  
Output a value of any printable type to the standard output (screen), and adds a newline
- `putChar :: Char -> IO ()`  
Writes the `Char` argument to the screen
- `getLine :: IO String`  
Read a line from the standard input and return it as a string
- The side effect of `getLine` is the consumption of a line of input, and the return value is a string
- `getChar :: IO Char`  
Read the next character from the standard input

Here are some more actions, `print`, its signature is `Show a => a -> IO ()`, this outputs a value of any printable type to the standard output, which is the screen and adds a new line. `PutChar` is a function from `Char -> IO ()`. It writes the `Char` argument that is provided to it on the screen. `getLine` is of type `IO string`, it reads the lines from the standard input and returns it as a `String`, the side effect of `getLine` is the consumption of line of input rather than the production of a line of output and the return value of `getLine` is a `String`. `getChar` is a function that reads the next character from the standard input.

(Refer Slide Time: 13:38)

## Binding

- `getLine` is of type `IO String`, but is there a way to use the return value?
- We need to **bind** the return value to an object of type `String` and use it elsewhere
- The syntax for binding is `<-`
- ```
main = do {  
    putStrLn "Please type your name!";  
    n <- getLine;  
    putStrLn ("Hello, " ++ n);  
}
```

We saw that `getLine` is of type `IO String`, but is there a way to use the value that is returned by `getLine`.

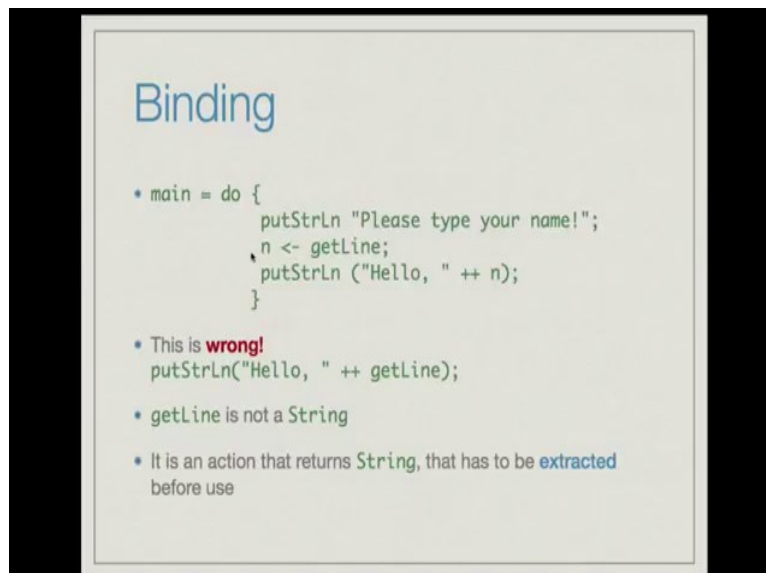
In other words we need to bind the return value of `getLine` to an object of type `String` and perhaps use it elsewhere. Haskell provides the following syntax for binding and the syntax is reminiscent of assignment in some other languages, which is just to say `<-`, it is like a left arrow for instance you could do this.

```
main = do {  
    putStrLn "Please enter your name!";  
    n <- getLine;  
    putStrLn ("Hello, " ++ n);  
}
```

`n` bound by `getLine` or in other words the output of `getLine` is bound to `n`.

And then I can use it here, in the following action, which is `putStrLn ("Hello, " ++ n)`; This has the effect of first asking the user for the name, waiting for the user to input her name and press the enter key and then printing Hello followed by her name on the screen.

(Refer Slide Time: 14:56)



The slide is titled "Binding" in blue. It contains a code snippet for a Haskell `main` function and a bulleted list of points.

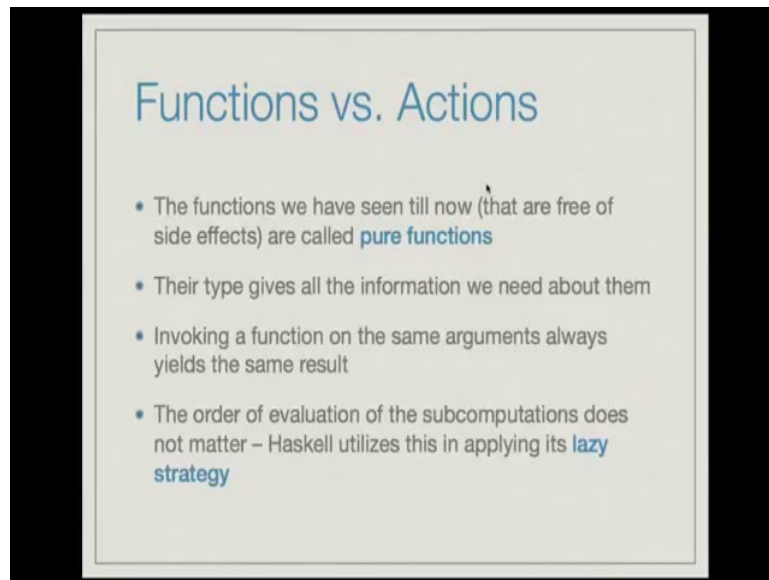
```
* main = do {  
    putStrLn "Please type your name!";  
    n <- getLine;  
    putStrLn ("Hello, " ++ n);  
}
```

- This is **wrong!**
`putStrLn("Hello, " ++ getLine);`
- `getLine` is not a `String`
- It is an action that returns `String`, that has to be **extracted** before use

Please note that this is wrong. `putStrLn ("Hello, " ++ getLine)`; this is because `++` is a so called pure Haskell function or operator and its arguments are list `a` and list `a` and the output is list `a`, the arguments are not of type `IO a`. Therefore, you cannot use `getLine` in the context

of a `++`, you should always bind it to some name and then use that name. `getLine` is not a `String`. It is an action that returns `String` that has to be extracted before the use. The extraction is the binding that happens here through `<-`.

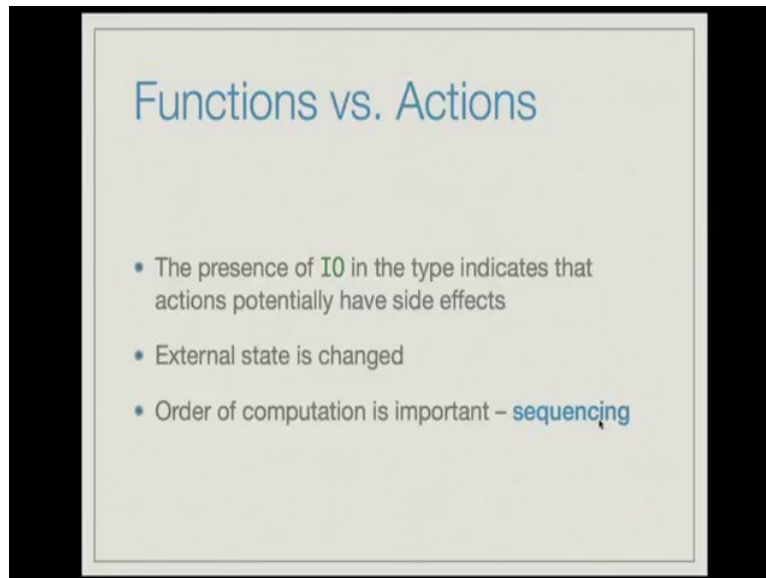
(Refer Slide Time: 15:47)



At this point we need to look a little closer at the distinction between functions and actions. A function that takes an integer as the argument and returns an integer as a result has type `Int -> Int`. An action that has a side effect in addition to consuming an integer and producing an integer has type `Int -> IO Int`. This distinction that Haskell maintains is in contrast to languages like C or Java, where the type signature of both functions that have side effects and functions that do not have side effects are just `Int -> Int`.

And in general any function is assumed to potentially produce a side effect, any function can produce a side effect there is nothing in the language itself that prevents functions from producing side effects. Haskell enforces this distinction between pure functions and actions. The functions that you have seen till now that are free of side effects are called pure function. Their type gives all the information we need about them. Invoking the function on the same arguments always yields the same result. And the order of evaluation of sub computations does not matter. Haskell utilizes this to great effect in applying its lazy strategy.

(Refer Slide Time: 17:20)

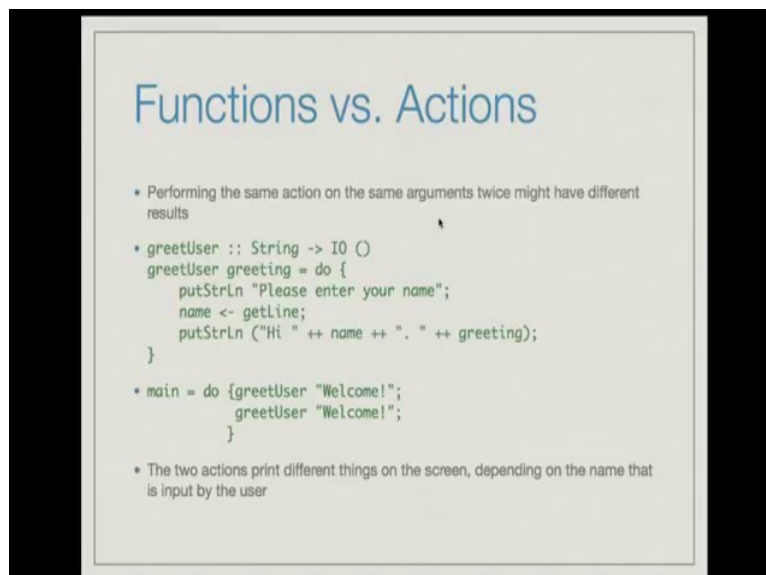


Functions vs. Actions

- The presence of **IO** in the type indicates that actions potentially have side effects
- External state is changed
- Order of computation is important – **sequencing**

In contrast to functions, actions usually have side effects and Haskell maintains this distinction by designating their output types with an IO. The presence of IO in the type indicates that actions potentially have side effects, external state is usually changed and the order of computation is important as in actions inside a do command. The actions take place in a sequence, so sequencing is something that is inherent to actions.

(Refer Slide Time: 18:03)



Functions vs. Actions

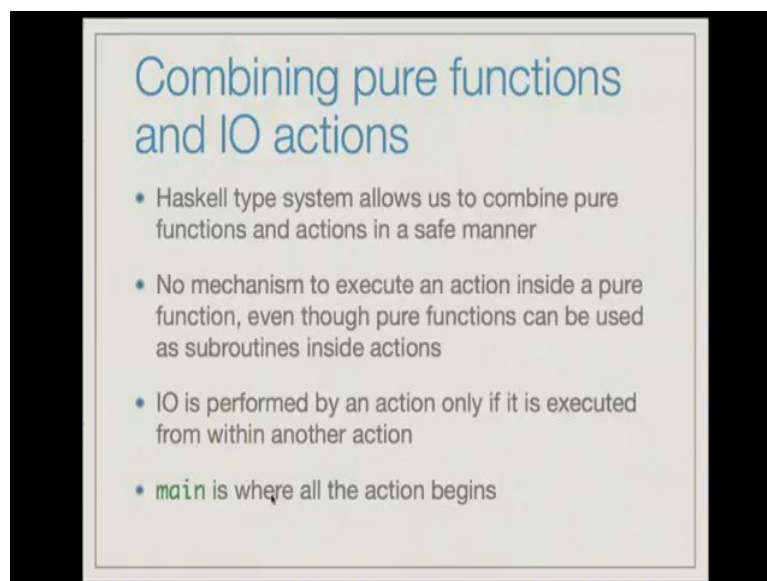
- Performing the same action on the same arguments twice might have different results
- ```
greetUser :: String -> IO ()
greetUser greeting = do {
 putStrLn "Please enter your name";
 name <- getLine;
 putStrLn ("Hi " ++ name ++ ". " ++ greeting);
}
```
- ```
main = do {greetUser "Welcome!";  
          greetUser "Welcome!";  
}
```
- The two actions print different things on the screen, depending on the name that is input by the user

Moreover, performing the same action on the same arguments twice might have different results. For instance consider the following action `greetUser`, which is of type `String -> IO ()`.

`greetUser greeting equals do. Inside the do block we have putStrLn “please enter your name“; which binds the value of getLine. The return value of get line is bound to name. ((ReferTime18:38)). Then, we put another string on the screen putStrLn (“Hi” ++ name ++ ”,” ++ greeting);`

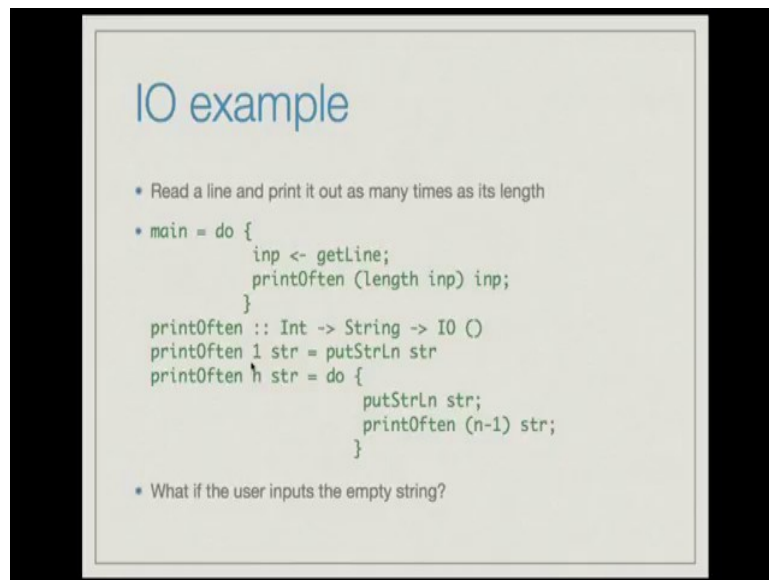
So, this greeting is something provided as an input. In main we can do the following `greetUser “Welcome!”; greetUser “Welcome”;`. You see that the `greetUser` function is being called with the exact same argument twice, but the two actions might print different things on this screen depending on the name that is input by the user at this point. Because, the `greetUser` action itself has a way to get input from the user and type the input back to the user namely the user's name. So, this shows that performing the same action on the same arguments twice might have different results and this is the fundamental difference between actions and functions.

(Refer Slide Time: 19:44)



One can combine pure functions and actions, but in a limited manner. we can use pure functions as subroutines and IO actions, but not the other way round. The Haskell type system allows us to combine pure functions and actions in a safe manner. No mechanism exists to execute an action inside a pure function even though pure functions can be used as subroutines inside actions. IO is performed by an action only if it is executed from within another action and main is, where all the action begins. So, main embed some actions inside it and each of those actions might be a do block with further actions inside.

(Refer Slide Time: 20:31)



Let us look at a few examples, here is one example, which is to read a line and print it out as many times as the length of the string that is input on the first line. Here is the program.

```
main = do {
    inp <- getLine;
    printOften (length inp) inp;
}

printOften :: Int -> String -> IO ()
printOften 1 str = putStrLn str
printOften n str = do {
    putStrLn str;
    printOften (n-1) str;
}
```

main equals do get a line and bind it to inp. inp stands for input. it is a variable that stores values of types String. Call the function printOften to print inp length inp times, if inp is of length n, then printOften input will be called and it will be printed n times.

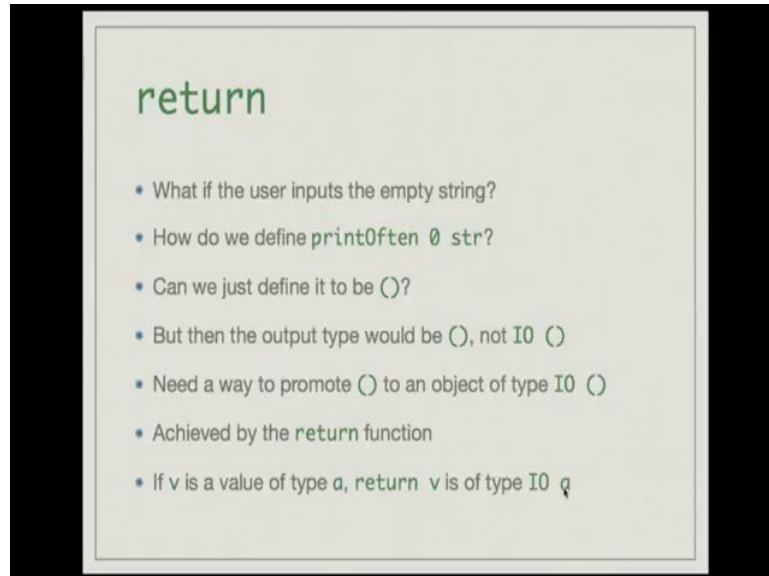
Thus achieved as follows. printOften is a function from Int -> String -> IO ().

printOften 1 str is just putStrLn str;

printOften n str is a recursive function. Inside a do block you first put putStrLn str; and then you printOften (n-1) str. So, this is an example of a non trivial

interaction with the user, but what if the user inputs the empty string. Notices that, if length of the input equals 0 there is no case here the catches it.

(Refer Slide Time: 22:03)



What if the user inputs the empty string, to handle this we need to define `printOften 0 str`. recall that the return value of `printOften` is `IO ()`. `printOften 0 str` - there should not be anything printed on the screen. So, can we just define it to be empty, but the output type then would be empty and not `IO empty`. So, you would get a type mismatch this means that we need a way to promote the empty tuple to an object of type `IO empty`. This is precisely achieved by the `return` function. The `return` function or the `return` action takes `v`, which is the value of type `a` and produces an action of type `IO a`, if `v` is a value of type `a`, `return v` is of type `IO a`.

(Refer Slide Time: 23:06)

IO example, fixed

- Read a line and print it out as many times as its length
- `main = do {
 inp <- getLine;
 printOften (length inp) inp;
}`

```
printOften :: Int -> String -> IO ()  
printOften 0 str = return ()  
printOften n str = do {  
    putStrLn str;  
    printOften (n-1) str;  
}
```

With this we can fix the earlier example as follows. `main` equals `do` get a line and bind it `input` `print` `often` `length` `input` `input` the crucial case is this `printOften 0 str = return ()`. This matches the type properly and it also handles this case in the case of not printing anything on the screen. Notice that here, there is no side effect as such, but still we designate this as type `IO` empty. So, this illustrates the fact that if an object is of type `IO` a it need not necessarily produce a side effect it only indicates the potential to produce a side effect whereas a pure function, which is not, whose type is not embedded to `IO` can never produce a side effect.

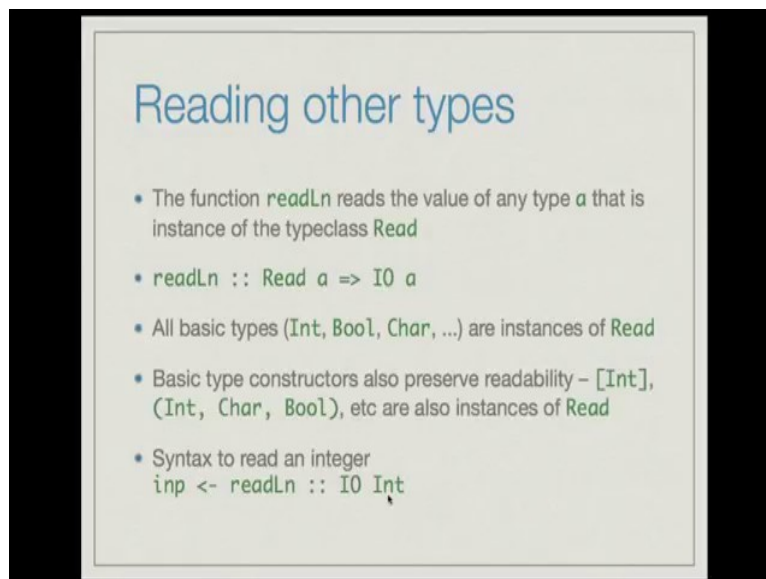
(Refer Slide Time: 24:13)

Another example

- Repeat an IO action n times
`ntimes :: Int -> IO () -> IO ()`
`ntimes 0 a = return ()`
`ntimes n a = do {
 a;
 ntimes (n-1) a;
}`
- Read and print 100 lines
`main = ntimes 100 act`
where
`act = do {
 inp <- getLine;
 putStrLn inp;
}`

Here is another example `ntimes`, which takes an integer and an action and repeat the action `n` times. This is the generalization of, what we did in the previous example, where we printed `inp` out to the screen some `n` times, where `n` is the length of the input. So, `ntimes 0 a` is just return empty, `ntimes n a` is do `a` followed by `ntimes (n-1) a`. `ntimes (n-1) a` which is to do the action `a` `(n-1)` times. Now, we can read and print 100 lines as follows `main = ntimes 100 act` where `act` equals this block, which reads an input from the user and prints out that input. `getLine` binded to `inp` and `putStrLn inp`. This action is repeated 100 times by `main`.

(Refer Slide Time: 25:15)



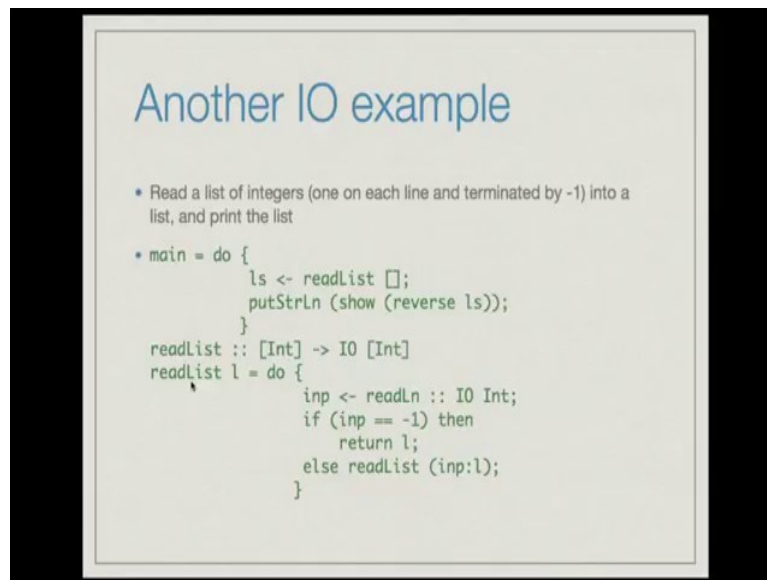
Reading other types

- The function `readLn` reads the value of any type `a` that is instance of the typeclass `Read`
- `readLn :: Read a => IO a`
- All basic types (`Int`, `Bool`, `Char`, ...) are instances of `Read`
- Basic type constructors also preserve readability – `[Int]`, `(Int, Char, Bool)`, etc are also instances of `Read`
- Syntax to read an integer
`inp <- readLn :: IO Int`

Strings are not the only things that can be read, we can read other values, but these other values have to belong to a type `a` that is an instance of the type class `Read`. For this we use the function `readLn`, which is denoted `readLn` whose type is `Read a => IO a`. All the basic types `Int`, `Bool`, `Char` etc are instances of `Read`. Therefore, you can use `readLn` to read Integers, Booleans Characters etc the basic type constructors also preserve readability.

So, for instance list `[Int]`, the triple `(Int, Char, Bool)` etc are also instances of `Read`. Here is the syntax to read an integer: `inp <- readLn :: IO Int`, and bind the result to `inp`. So, since `readLn` is supposed to cater to the reading of any of these types that belongs to class `Read` you need to specify, which type you want the input to be.

(Refer Slide Time: 26:38)



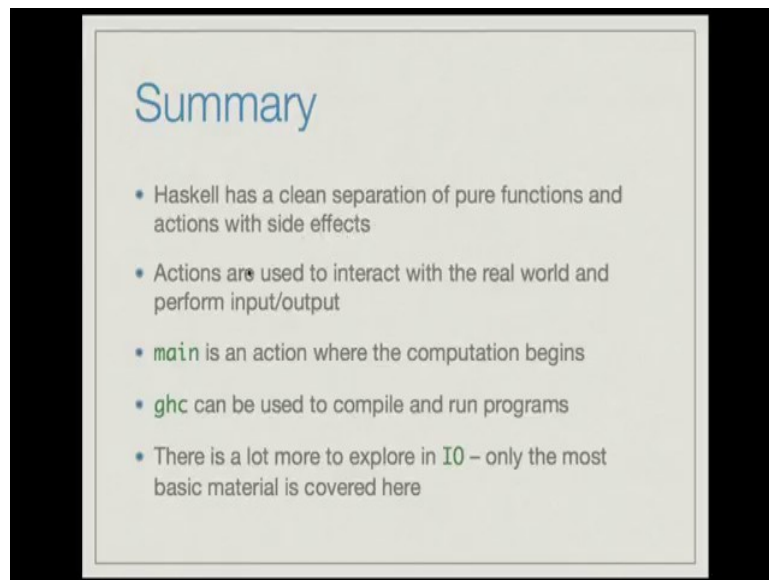
So, here is an example, in this example we read a list of integers one on each line and finally, terminated by -1 into a list and print the list.

```
main = do {  
    ls <- readList [];  
    putStrLn (show (reverse ls));  
}
```

The return value of readList is empty list. So, readList is a function you are providing the empty list as an input to read list. It is of type [Int] -> IO [Int]. So, it produces some side effects and its return value is a list of integers, which is what is bound to the ls here and then it so happens that readList reads the list in reverse order.

So, we show the reverse of ls and invoke putStrLn on it, in this way we read a bunch of integers one on each line and output them as a list. Here is the description of read list. read list l equals do readLn as an integer, this you denote by saying inp <- readLn :: IO Int and bind the return value to inp. If inp is -1, then, there is nothing more to do, so you just return l, which is the list that was provided as input else you have read the input you add the input to the list by saying inp:l you add it to head of the list and you proceed with the readList function presumably reading more input.

(Refer Slide Time: 28:27)



In summary Haskell has a clean separation of pure functions and actions with side effects, actions are used to interact with the real world and perform input output .main is the action where the computation begins, ghc can be used to compile and run programs. So, there is a lot more to explore in IO but only the most basic material is covered here which suffices for rudimentary interaction with users.